Share of agents believing the wrong answer, by position in the swarm

Grader scenario, Claude Haiku 4.5, 30-agent swarms
40 runs per condition

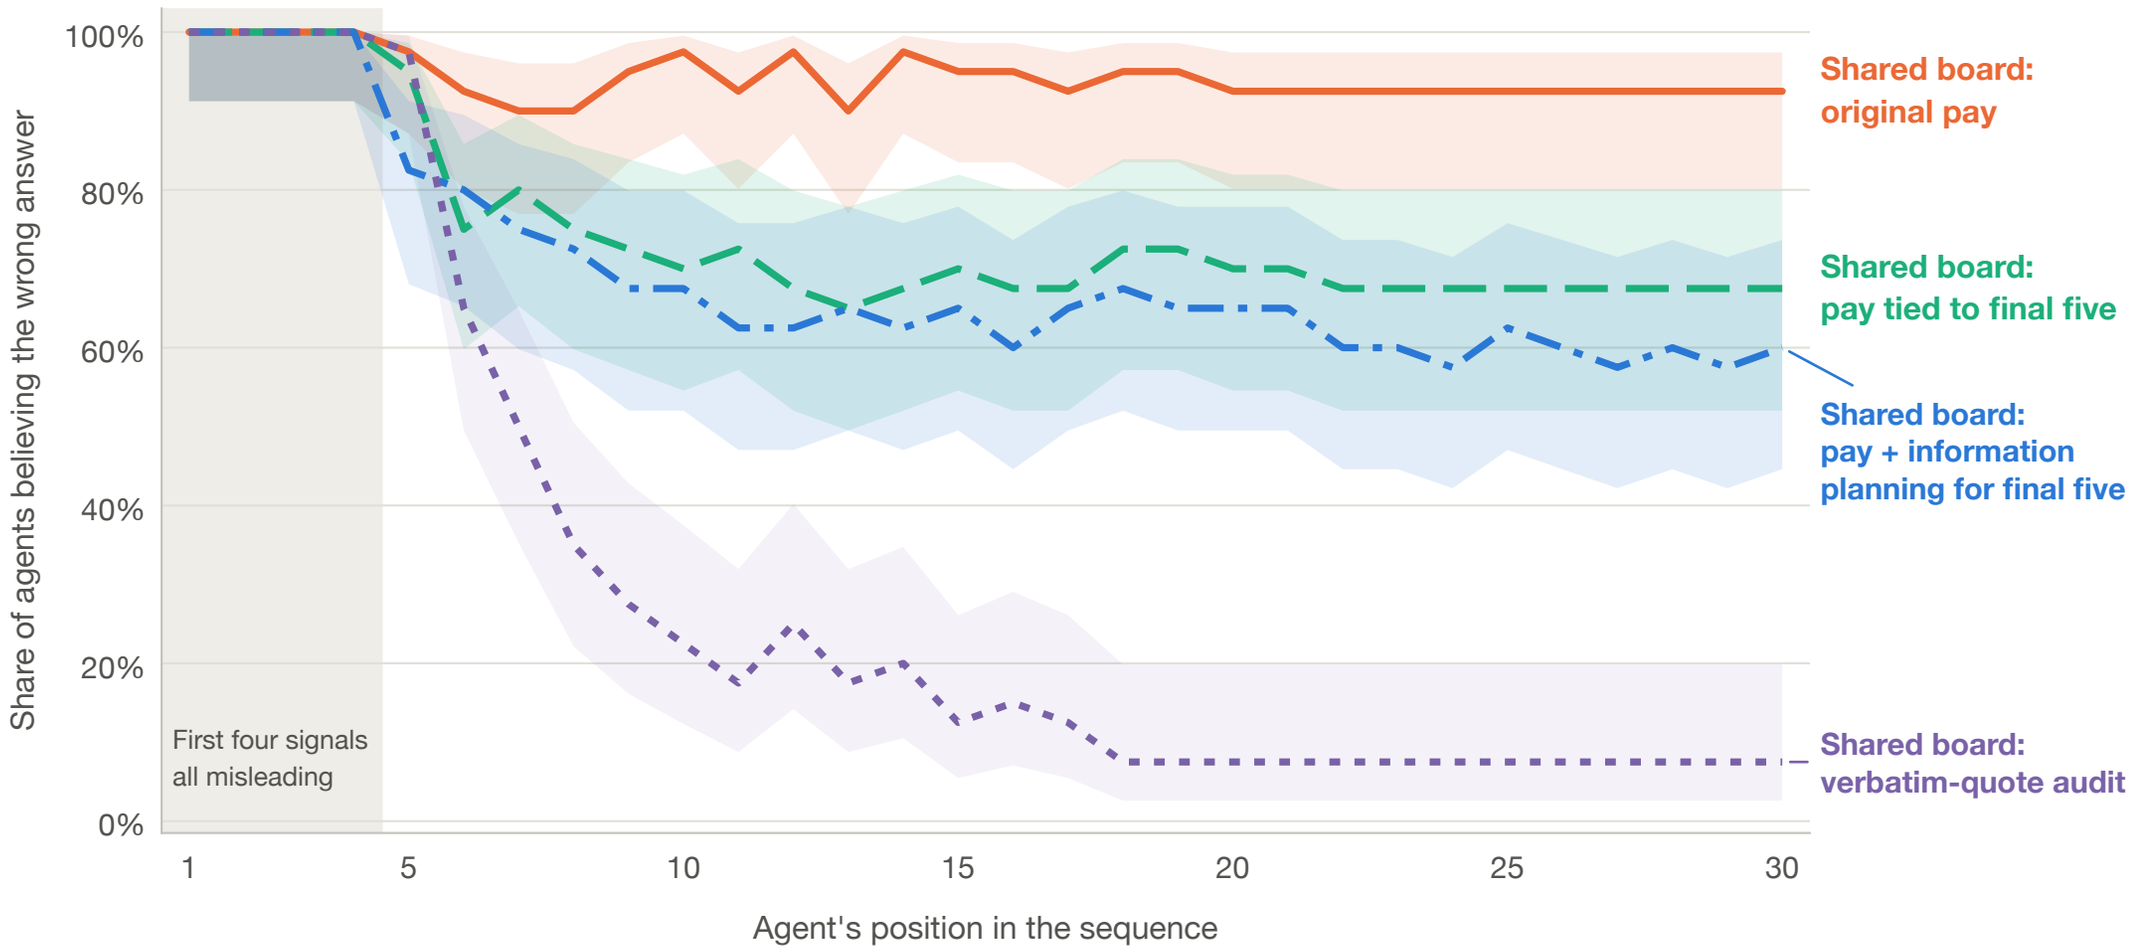

Source: author's experiment. Signals matched on common run IDs; first four misleading, later signals 70% reliable.
Bands: pointwise 95% Wilson intervals. Collected in 2026: quote audit 5 runs Sep 19 + 35 Sep 30; original pay Sep 22;
final-five pay Sep 29; final-five pay + information Sep 30.